

The BPA Lab

**Architecture and Documentation of a Demonstration Factory for Business Process
Automation and Analytics**

Matthias Zapp

May 23, 2026

Table of contents

1. Overview	1
1.1. About this book	1
1.2. Who this book is for	2
1.3. Source code and related repositories	2
1.4. Status and roadmap	2
1.5. Licence and citation	3
2. Introduction	5
2.1. What is the BPA Lab?	5
2.2. Two goals: demonstration and teaching	5
2.2.1. Demonstration	5
2.2.2. Teaching	6
2.3. Why a demonstration factory?	6
2.4. A bird's-eye view	6
3. Solution Overview	7
3.1. Three architectural layers	7
3.1.1. Layer 1 — Controllers	7
3.1.2. Layer 2 — Process applications (job workers)	7
3.1.3. Layer 3 — BPMS workflow engine	8
3.2. A modular, domain-driven design	8
3.3. A simplified bird's-eye view	8
3.4. The demonstration factory implementation	9
4. Business Scenario of the Demonstration Factory	11
4.1. The scenario in a nutshell	11
4.2. Order management	11
4.3. Production control	13
4.4. Purchasing	13
4.5. Manufacturing	13
4.6. Shipment	15
4.7. Warehouse operations	15
4.8. Teaching note	16
5. Data and Data Sources	17
5.1. Overview of data sources	17
5.2. The MySQL data model	17
5.3. Sample and master data	18
5.4. Process data and the foundation for process mining	18
5.5. Data lifecycle and analytics	19

6. Software Architecture	21
6.1. Requirements	21
6.1.1. R1 — Independent operation of components	21
6.1.2. R2 — Integration of external IT systems	21
6.1.3. R3 — Clear and distinct responsibilities	21
6.1.4. R4 — Easy extensibility	22
6.1.5. R5 — Interchangeability and independent evolution	22
6.1.6. R6 — Comprehensive and flexible process data capture	22
6.1.7. R7 — User-friendliness (operator perspective)	22
6.1.8. R8 — Easy installation and low infrastructure requirements (developer perspective)	22
6.1.9. R9 — Robustness	22
6.1.10. R10 — Maintainability and operation	22
6.2. Architecture in the C4 model	22
6.2.1. Context diagram	22
6.2.2. Container diagram	23
6.3. Mapping requirements to architecture	23
6.4. Going deeper	24
7. Architecture Decisions	25
7.1. ADR-001 — MQTT broker for job-worker-controller communication	25
7.2. ADR-002 — Message Send/Receive Tasks instead of Events for inter-process communication	25
7.3. ADR-003 — Messaging between processes and process applications	26
7.4. ADR-004 — Camunda 8 self-managed with Docker Compose (Camunda 8 Core)	26
7.5. ADR-005 — Job worker for sending emails	26
7.6. ADR-006 — Job worker for database transactions; MySQL in a Docker container	27
7.7. Reflection questions	27
Appendices	29
A. Components and Sister Repositories	29
A.1. Main repositories	29
A.1.1. bpa_lab_docs	29
A.1.2. bpa_lab_demonstration_factory	29
A.1.3. bpa_lab_student_docs	30
A.2. Hardware components	30
A.3. Software components	30
A.4. Further reading	30
B. Glossary	31
C. References	33

1. Overview




Figure 1.1.: BPA Lab — Conceptual Architecture

This book documents the **Business Process Automation Lab (BPA Lab)** at TH Köln — a small-scale, modular model factory that demonstrates and teaches concepts and technologies for business process automation and analytics.

1.1. About this book

The BPA Lab combines **software components** (Business Process Management Systems, Robotic Process Automation, Process Mining) with **hardware components** (a Fischertechnik Industry 4.0 factory, a Fischertechnik warehouse robot, and IoT devices) to simulate the production and logistics processes of a bicycle manufacturer.

This book aims to provide:

1. A **high-level overview** of the BPA Lab and its goals.
2. A **description of the business scenario** — the end-to-end ordering, manufacturing, and shipment process implemented in the lab.
3. A **description of the software architecture** and the design decisions taken along the way.

It is structured as follows:

1. Overview

1. **Introduction:** Motivation, goals, and scope of the BPA Lab.
2. **Solution overview:** The three architectural layers — controllers, process applications, workflow engine.
3. **Business scenario:** The end-to-end bicycle ordering, manufacturing, and shipment process, broken down into its sub-processes (order management, production control, purchasing, manufacturing, shipment, warehouse operations).
4. **Data and data sources:** Data model, data flows, and persistence in the BPA Lab.
5. **Software architecture:** Requirements and a C4 view of the system.
6. **Architecture decisions:** Documented decisions and their rationales.

The appendices contain:

- **Appendix A:** Overview of related components and sister repositories.
- **Appendix B:** A glossary of key terms.
- **References:** All cited literature.

1.2. Who this book is for

This book is written for:

- **Lecturers and researchers** who want to teach business process automation with a tangible, end-to-end reference system.
- **Students** working on bachelor's or master's projects inside the BPA Lab.
- **Practitioners** looking for a concrete reference architecture combining a BPMS, IoT hardware, and process mining.

1.3. Source code and related repositories

The implementation of the demonstration factory is maintained as a separate repository: [bpa_lab_demonstration_factory](#).

Documentation for student component projects is available in the wiki of [bpa_lab_student_docs](#).

1.4. Status and roadmap

! Work in progress

The BPA Lab — and this book — are under continuous development. Several sections are marked as *work in progress* and will be expanded in future revisions. Bug reports and suggestions are very welcome via the [issue tracker](#).

1.5. Licence and citation

This work is published under the **Creative Commons Attribution-ShareAlike 4.0** licence (CC BY-SA 4.0).

Suggested citation:

Zapp, Matthias (2026). *The BPA Lab — Architecture and Documentation of a Demonstration Factory for Business Process Automation and Analytics*. TH Köln (University of Applied Sciences). Online: https://bpalabthcologne.github.io/bpa_lab_book/.

2. Introduction

Learning objectives

After reading this chapter, you will be able to:

- describe the purpose and main goals of the BPA Lab,
- distinguish between its *demonstration* and *teaching* missions,
- identify the main hardware and software components and their roles.

2.1. What is the BPA Lab?

The **Business Process Automation Lab (BPA Lab)** at TH Köln is a small, modular model factory dedicated to business process automation and analytics. It combines real-world software components used in industry — Business Process Management Systems (BPMS), Robotic Process Automation (RPA), and process-mining tools — with physical hardware: a Fischertechnik Industry 4.0 factory, a Fischertechnik warehouse robot, and a range of IoT devices.

Together, these components simulate the **production and logistics processes of a bicycle manufacturer** — from the customer order to the delivery of a custom-built bike.

2.2. Two goals: demonstration and teaching

The BPA Lab pursues two complementary goals.

2.2.1. Demonstration

The lab demonstrates concepts and technologies for business process automation and analytics in a tangible way. It shows how:

- a central BPMS (Camunda 8) orchestrates end-to-end business processes,
- modular **job workers** implement the actual work and interact with hardware,
- **process mining** can be used to analyse process executions and identify improvement opportunities.

By bringing together software and hardware in one running system, the lab makes abstract concepts immediately visible.

2. Introduction

2.2.2. Teaching

Individual building blocks of the BPA Lab — for instance the Fischertechnik warehouse robot — are also used in **student projects** (bachelor's and master's level) as starting points for self-defined process implementations. This enables a hands-on form of teaching for **process-oriented application integration**: students do not merely model processes on paper but see them execute in a real, physical system.

2.3. Why a demonstration factory?

Business process management theory and notations like BPMN are often taught using abstract examples — pizza ordering, expense approvals, generic procurement workflows. These examples are accessible but rarely convey:

- the **technical complexity** of integrating heterogeneous systems,
- the **operational concerns** of running long-lived process instances,
- the **data perspective** that is at the heart of process mining and analytics.

A demonstration factory bridges this gap. Students and visitors can observe how a customer order arrives via a web form, triggers a workflow in Camunda 8, dispatches commands via MQTT to physical hardware, and eventually leads to a real, manufactured product being moved through a real, physical warehouse.

2.4. A bird's-eye view

The next chapter — [Solution overview](#) — introduces the three architectural layers of the BPA Lab. Subsequent chapters describe the business scenario, the data perspective, the software architecture, and the architecture decisions taken so far.

How to read this book

If you are new to the BPA Lab, read the chapters in order. If you are already familiar with the lab and want to consult a specific aspect — for example a particular sub-process or an architectural decision — feel free to jump directly to the relevant chapter using the sidebar or the search box at the top of the page.

3. Solution Overview

Learning objectives

After reading this chapter, you will be able to:

- name the three architectural layers of the BPA Lab and explain their responsibilities,
- describe the role of Camunda 8 as the central workflow engine,
- explain how hardware components and process applications interact,
- justify the modular, domain-driven design approach used in the BPA Lab.

3.1. Three architectural layers

The architecture and implementation of the BPA Lab are under continuous development. The solution is structured into three clearly separated layers.

3.1.1. Layer 1 — Controllers

Controllers are responsible for the direct control of hardware components: the Fischertechnik Industry 4.0 factory, the Fischertechnik warehouse robots, and IoT devices. They are predominantly implemented in **Python** and translate high-level commands (for example “*transport workpiece from slot A to slot B*”) into the low-level signals that move motors and read sensors.

3.1.2. Layer 2 — Process applications (job workers)

Process applications are middleware components that implement the actual business logic of individual work steps. In Camunda 8 terms, they are **job workers**: they subscribe to specific service tasks and execute them when activated by the workflow engine.

Job workers communicate:

- with the **workflow engine** via **gRPC** (the standard Camunda 8 / Zeebe protocol),
- with the **controllers** via **MQTT** (a lightweight publish-subscribe protocol that is well established in the IoT space).

Process applications are predominantly implemented in **Java**.

3. Solution Overview

3.1.3. Layer 3 — BPMS workflow engine

The central workflow engine is **Camunda 8 (self-managed)**. It executes:

- **BPMN** process models — describing how work flows through the system,
- **DMN** decision models — encoding business rules.

Camunda 8 maintains the state of each running process instance, handles incidents, and provides the data foundation for process mining and analytics.

3.2. A modular, domain-driven design

A key design principle in the BPA Lab is **modularity**. The system follows a **Domain-Driven Design (DDD)** approach so that:

- **multiple students and developers can work in parallel** on different parts of the system without stepping on each other's toes,
- **components can be reused** in student projects to build new process applications,
- **components can be replaced or upgraded** independently.

The primary criterion for splitting job workers and process models is the **business domain**:

- Order management
- Manufacturing
- Shipping
- Purchasing
- Logistics / warehouse operations

This separation aligns with the bounded contexts familiar from DDD and makes the system easier to reason about, both for teachers explaining it and for students extending it.

3.3. A simplified bird's-eye view

The following diagram summarises the interaction between the three layers and the hardware they ultimately drive.

The customer-facing entry point is typically a **web application** for order capture. From there, work flows downwards through Camunda 8 and the job workers and eventually reaches the physical Fischertechnik components. Events generated along the way flow back upwards and are persisted in databases for later analysis.




Figure 3.1.: Simplified architecture of the BPA Lab

3.4. The demonstration factory implementation

The latest version of the BPA Lab as a runnable **demonstration factory** is available in a separate repository:

github.com/BpaLabTHCologne/bpa_lab_demonstration_factory

The implementation is fully containerised with **Docker**. This makes local setup straightforward and ensures that the environment is reproducible across different machines — a key requirement when students set up the system on their own laptops.

💡 Getting started in practice

If you are setting up the BPA Lab for the first time, start with the `docker-compose.yml` of the demonstration factory and follow the installation guide there. To develop your own components on top of the lab, the [student-docs wiki](#) provides component-level guides.

4. Business Scenario of the Demonstration Factory

Learning objectives

After reading this chapter, you will be able to:

- describe the end-to-end business processes of a custom bicycle manufacturer in BPMN terms,
- explain the orchestrating role of the order management process,
- distinguish between the responsibilities of the order management, production control, purchasing, manufacturing, shipment, and warehouse operations processes,
- recognise how Fischertechnik hardware is used to *physically* execute parts of these processes.

Status of this chapter

This chapter describes the business processes of the BPA Lab at a strategic level. It is *work in progress* — terminology, level of detail, and harmonisation of process models are being refined as the lab evolves.

4.1. The scenario in a nutshell

A bicycle manufacturer offers highly customisable bicycles to consumers.

The entire end-to-end process is orchestrated by the **order management process**, which in turn initiates purchasing, manufacturing, and shipping processes as needed.

The manufacturing processes use the **Fischertechnik Industry 4.0 factory** to post goods receipts of components and to perform the actual manufacturing steps. Order management and shipping additionally use the **distribution warehouse** (a separate Fischertechnik robot) to post goods receipts and goods issues for finished products.

The following sections describe each sub-process in turn. The process models do **not** include technical aspects and show only the **happy path** of execution.

4.2. Order management

The order management process governs the fulfilment of customer orders.

4. Business Scenario of the Demonstration Factory




Figure 4.1.: Process landscape of the BPA Lab




Figure 4.2.: Order management process

It starts when a customer order is entered into a web form, which is pre-filled with product data (bicycle models) from the database. During execution, the **availability of the ordered products** is checked.

Two main scenarios apply:

- **All products in stock:** the shipping process is triggered directly. Once shipping completes successfully, the process is finished.
- **One or more products out of stock:** for each unavailable product, the production control process is triggered, which creates a **production order**. The production order specifies the required components.
 - If the components are also unavailable, a **purchasing process** is started first.
 - Otherwise, the **manufacturing process** is triggered immediately.

Once the missing products are manufactured, shipping is initiated.

4.3. Production control




Figure 4.3.: Production control process

The production control process manages the creation and execution of **one production order for one product**. It determines the component demand and — if needed — triggers a purchasing process to procure the missing components. Once components are available, it triggers the actual manufacturing process.

4.4. Purchasing

The purchasing process manages the procurement of required components. It starts whenever a component is needed and creates a **purchase order**.

The purchase order can be adjusted by the user — for example by selecting a specific vendor. Materials are then stored and stock levels are updated. Upon user confirmation (or after a configurable timeout), the stock level is increased in the inventory database.

4.5. Manufacturing

The manufacturing process governs the **physical manufacturing simulation** on the Fischertechnik factory for **one item of one product type**.

4. Business Scenario of the Demonstration Factory




Figure 4.4.: Purchasing process




Figure 4.5.: Manufacturing process

It starts after a production order has been created by the production control process and sufficient components are available.

The process must place a manufacturing order that produces a new bicycle from the specified material. Before that, however, it must check whether a **physical workpiece** is already available in the warehouse of the Fischertechnik factory. If not, a physical workpiece must first be transferred to the factory so that the control program can store and recognise it. Once the bicycle has been produced from the material (workpiece), the process is complete.

4.6. Shipment




Figure 4.6.: Shipment process

The shipment process manages the dispatch of finished goods to the customer. It is **only** triggered if a product is in stock and ready to be shipped.

The flow is:

1. Shipping data is checked.
2. The distance and duration from the warehouse to the delivery address are computed — the **Openrouteservice API** is used for this.
3. The logistics sub-process retrieves the customer order from the warehouse.
4. A message is sent to the customer to inform them about the dispatch.
5. Upon handover of the order to the customer, both the shipping process and the surrounding end-to-end process instance are completed.

4.7. Warehouse operations

The warehouse operations process governs the flow of finished goods in a distribution warehouse — represented physically by the Fischertechnik warehouse robot.

It starts when a **transfer order** for the warehouse is received from another business process. A transfer order can be either an instruction to **store products** in the warehouse or to **retrieve products** from it.

4.8. Teaching note

Process landscape as a teaching tool

This process landscape works well to introduce key BPMN concepts in context:

- **call activities** (e.g. order management calling production control),
- **message-based communication** between long-running processes,
- **boundary events** to handle exceptions raised by job workers,
- **multi-instance activities** for handling several products in a single order.

By stepping through the running system together with students, abstract notation suddenly becomes concrete.

5. Data and Data Sources

Learning objectives

After reading this chapter, you will be able to:

- identify the main data sources in the BPA Lab,
- describe the role of the relational database in the overall architecture,
- explain how process and event data become the foundation for process mining.

5.1. Overview of data sources

The BPA Lab integrates data from several sources — from **master data** in a relational database, via **process state** maintained by the workflow engine, to **sensor and telemetry data** from the Fischertechnik components.

The following figure gives an overview of the main data sources and the kinds of data they hold:




Figure 5.1.: Data overview in the BPA Lab

5.2. The MySQL data model

The following diagram shows the concrete data model of the MySQL database, including the entities and attributes of the current implementation:

The database serves as the central store for:

- **Product master data** — bicycle models and their components,
- **Customer data** and delivery addresses,
- **Stock levels** for components and finished goods,

5. Data and Data Sources




Figure 5.2.: MySQL data model

- **Orders** — customer orders, production orders, purchase orders — together with their status history,
- **Vendor data.**

5.3. Sample and master data

Documentation and sample data of the Fischertechnik factory are available in the [Sample_data](#) folder of the documentation repository. This data serves as:

- **initial master data** when setting up a new instance of the BPA Lab,
- **test data** for developing new components,
- a **demonstration baseline** for teaching sessions and presentations.

5.4. Process data and the foundation for process mining

In addition to persistent master data, the running processes generate substantial amounts of **process data**:

- **Camunda 8** stores process instance data, variables, tokens, and lifecycle events.
- **MQTT messages** record the communication between job workers and controllers.
- **Job worker logs** capture technical operations and exceptions.

This event data is the foundation for **process mining** — a central teaching and demonstration component of the BPA Lab. Process mining allows students to:

- **discover** the actual flow of processes from event logs (and compare them to the modelled flow),
- **check conformance** between modelled and actual behaviour,
- **analyse performance** — durations, bottlenecks, rework loops.

 Extension idea

For process-mining exercises, it is worth writing an event-log export script that converts Camunda 8 data into the **XES** (eXtensible Event Stream) format. This enables direct integration with tools such as ProM, Disco, or PM4Py, and gives students a realistic data pipeline from running process to analytical tool.

5.5. Data lifecycle and analytics

Looking at the data lifecycle as a whole:

1. **Master data** is loaded once when the lab is set up and updated when new products, customers, or vendors are added.
2. **Operational data** — orders, stock movements, production orders — is created continuously as processes run.
3. **Event data** is emitted by the workflow engine, the job workers, and the MQTT broker as side effects of process execution.
4. **Analytical data** is derived from the above by export and transformation pipelines, then fed into process-mining and BI tools.

This layered view of data is itself a teaching opportunity: it mirrors the classical separation of OLTP and OLAP systems that students will encounter in industry.

6. Software Architecture

Learning objectives

After reading this chapter, you will be able to:

- name the high-level requirements that shaped the BPA Lab architecture,
- describe the architecture using the C4 model,
- relate individual architectural choices to the requirements that motivated them.

Status

The architecture is under continuous development. Several aspects are still being discussed. This chapter captures the current state.

6.1. Requirements

The following high-level requirements were gathered at the very beginning of the BPA Lab project (originally as part of a thesis project). They served as the basis for the first software architecture and remain the reference point against which architectural choices are evaluated.

6.1.1. R1 — Independent operation of components

The individual components and processes should be able to run independently of each other. Each process must be startable on its own, without depending on a triggering instance from a higher layer.

6.1.2. R2 — Integration of external IT systems

The system must allow integration with existing IT components — for example with the web application used for order capture, or with tools that flexibly extract data for process mining and analytics.

6.1.3. R3 — Clear and distinct responsibilities

The responsibilities of individual components must be clearly defined and demarcated from one another, so that the system remains understandable as it grows.

6. Software Architecture

6.1.4. R4 — Easy extensibility

The system must be easy to extend with additional, self-developed components.

6.1.5. R5 — Interchangeability and independent evolution

Components must be developable, deployable, and replaceable independently of one another.

6.1.6. R6 — Comprehensive and flexible process data capture

Process data must be comprehensively capturable. It must be stored and made flexibly available for a range of analytical use cases — in particular process mining.

6.1.7. R7 — User-friendliness (operator perspective)

The user interface must be intuitive and easy to understand, ensuring a positive user experience for those interacting with the lab during demonstrations and student sessions.

6.1.8. R8 — Easy installation and low infrastructure requirements (developer perspective)

The system must be easy to install. Students must be able to host their own instances on standard laptops without specialised infrastructure.

6.1.9. R9 — Robustness

The architecture must be robust. Failures must be minimised and reliable functionality must be ensured under a wide range of operating conditions.

6.1.10. R10 — Maintainability and operation

The system must be easy to maintain over its lifetime.

6.2. Architecture in the C4 model

The architecture is documented using the **C4 model** (Brown 2018) — a notation for describing software systems at different levels of abstraction: Context, Container, Component, and Code. Note that *containers* in the C4 sense are **not** the same as Docker containers.

6.2.1. Context diagram

The context diagram shows the BPA Lab in relation to its users and external systems.

6.3. Mapping requirements to architecture




Figure 6.1.: C4 context diagram

6.2.2. Container diagram

The container diagram shows the main technical building blocks of the BPA Lab and how they interact.

6.3. Mapping requirements to architecture

The following table illustrates how individual requirements are addressed by architectural choices. It is not exhaustive — but it shows the line of reasoning from requirement to design.

Requirement	Realisation in the architecture
R1 — Independent operation	Loose coupling via MQTT and gRPC; self-contained job workers
R2 — Integration of external systems	Standard protocols (REST, MQTT, gRPC); Camunda connector concept
R3 — Clear responsibilities	Domain-driven separation of job workers and process models
R6 — Process data capture	Camunda 8 as central event hub; MySQL for operational data
R8 — Easy installation	Docker Compose setup; self-managed Camunda 8 Core
R5/R9 — Interchangeability / robustness	Asynchronous messaging; clear interfaces between layers

6. Software Architecture




Figure 6.2.: C4 container diagram

6.4. Going deeper

💡 Further reading

The original sources on the C4 model — including examples, decision records, and tools like [Structurizr](https://c4model.com/) — are available at <https://c4model.com/>. For a deeper dive into Domain-Driven Design as the principle behind the modular separation, see Evans (Evans 2003).

7. Architecture Decisions

Learning objectives

After reading this chapter, you will be able to:

- explain the rationale behind the central architectural choices in the BPA Lab,
- read and write lightweight Architecture Decision Records (ADRs),
- reflect on typical trade-offs in BPM-centric architectures.

This chapter records selected architecture questions raised during the BPA Lab project, together with the corresponding decisions and justifications. The list is **not exhaustive** and is updated as the lab evolves.

Each decision follows a lightweight **ADR** (Architecture Decision Record) schema: context, decision, and rationale are captured succinctly.

7.1. ADR-001 — MQTT broker for job-worker–controller communication

Status: Decided.

Decision: A MQTT broker is used for communication between process applications (job workers) and controllers.

Rationale:

- decoupling of senders and receivers through the publish-subscribe pattern,
- MQTT is a well-established standard protocol in the IoT space,
- lightweight footprint, well suited to the hardware components in use.

7.2. ADR-002 — Message Send/Receive Tasks instead of Events for inter-process communication

Status: Decided.

Decision: Inter-process communication is modelled in BPMN using **Message Send Tasks** and **Message Receive Tasks**, not events.

Rationale:

- Tasks allow the use of **boundary events** — for example to handle errors raised by job workers,
- complex error paths and timeouts are easier to express.

7.3. ADR-003 — Messaging between processes and process applications

Status: Decided.

Decision: Communication between processes and process applications is implemented using messages as well.

Rationale:

- a simple, uniform solution,
- consistent with the communication pattern between process applications and process control.

7.4. ADR-004 — Camunda 8 self-managed with Docker Compose (Camunda 8 Core)

Status: Decided.

Decision: The self-managed variant of Camunda 8 is used in combination with Docker Compose and Camunda 8 Core.

Rationale:

- avoids interferences between contributors during the development phase,
- development and production-like environments are identical,
- Kubernetes is intentionally not used — the BPA Lab is not a production environment in the classical sense.

Open question: Performance with Docker and the resulting number of containers needs to be verified as the system grows.

7.5. ADR-005 — Job worker for sending emails

Status: Decided (reversible).

Decision: Email sending is implemented through a dedicated job worker rather than via the SendGrid connector.

Rationale:

- issues encountered with the SendGrid connector at the time of the decision,
- an existing solution was already in place in the project,
- to be re-evaluated as Camunda 8 connectors continue to evolve.

7.6. ADR-006 — Job worker for database transactions; MySQL in a Docker container

Status: Decided.

Decision: Database transactions are handled by a dedicated job worker; the MySQL database runs in a Docker container.

Rationale:

- existing solution available in the project,
- greater flexibility than the connectors available at the time of the decision.

7.7. Reflection questions

 For your own consideration

1. ADR-001 selects MQTT. What **alternative protocols** could have been considered, and what trade-offs would they imply?
2. ADR-005 is explicitly marked as reversible. Which **criteria** would you use to decide whether to switch back to a connector-based solution?
3. ADR-004 deliberately avoids Kubernetes. Under which circumstances would that decision need to be revisited?

A. Components and Sister Repositories

This appendix gives an overview of the repositories that together make up the BPA Lab ecosystem, and lists the main hardware and software components.

A.1. Main repositories

A.1.1. bpa_lab_docs

https://github.com/BpaLabTHCologne/bpa_lab_docs

Role: central documentation and architecture of the BPA Lab as a demonstration factory. The source of the content in this book.

Contains:

- overarching architecture diagrams,
- sample data of the Fischertechnik factory,
- BPMN process models (overview),
- C4 architecture diagrams,
- updated graphics.

A.1.2. bpa_lab_demonstration_factory

https://github.com/BpaLabTHCologne/bpa_lab_demonstration_factory

Role: current implementation of the demonstration factory. A complete, runnable setup based on Docker.

Contains:

- self-managed Camunda 8 setup,
- all job worker implementations (Java),
- BPMN and DMN models of the business processes,
- controller implementations (Python),
- `docker-compose.yml` for local deployment.

A. Components and Sister Repositories

A.1.3. bpa_lab_student_docs

https://github.com/BpaLabTHCologne/bpa_lab_student_docs/wiki

Role: documentation for student projects. Describes how individual components (e.g. the warehouse robot) can be used in self-developed process implementations.

Contains:

- per-component how-to guides,
- example projects,
- notes on setting up subsystems independently.

A.2. Hardware components

Component	Purpose	Control
Fischertechnik Industry 4.0 factory	Goods receipt and manufacturing	Python controller, MQTT
Fischertechnik warehouse robot	Storage and retrieval in the distribution warehouse	Python controller, MQTT
IoT devices	Sensor data and supplementary telemetry	Python controller, MQTT

A.3. Software components

Component	Technology	Main responsibility
Camunda 8 (self-managed)	BPMS, based on Zeebe	Workflow engine and process orchestration
Job workers	Java, gRPC	Execution of service tasks
Controllers	Python	Hardware control
MQTT broker	e.g. Mosquitto	Asynchronous communication
MySQL	Relational database	Master data and stock data
Web application	Frontend	Order capture and user interaction
Openrouteservice	External API	Distance and routing calculation in the shipment process

A.4. Further reading

As the BPA Lab is under continuous development, it is always worth checking the respective repository READMEs for the most current information.

B. Glossary

This appendix explains the technical terms used throughout the book.

- ADR (Architecture Decision Record)** Lightweight documentation format for recording architectural decisions in a structured way — context, decision, and rationale.
- BPA (Business Process Automation)** The use of technology to automate the execution of business processes.
- BPMN (Business Process Model and Notation)** A standard of the Object Management Group (OMG) for graphical modelling of business processes. Current version: 2.0.2 (Object Management Group 2014).
- BPMS (Business Process Management System)** A platform for modelling, executing, monitoring, and optimising business processes. The BPA Lab uses Camunda 8.
- Camunda 8** A cloud-native BPMS built around the Zeebe workflow engine. The BPA Lab uses the *self-managed* variant (Camunda Services GmbH 2026).
- Controller** In the BPA Lab, a software component that directly controls hardware (Fischertechnik components, IoT devices). Predominantly implemented in Python.
- C4 model** A notation for visualising software architecture at four levels of abstraction: Context, Container, Component, Code. Introduced by Simon Brown — <https://c4model.com/> (Brown 2018).
- DDD (Domain-Driven Design)** An approach to software development that places the business domain at the centre. The guiding principle for the modular decomposition of the BPA Lab (Evans 2003).
- DMN (Decision Model and Notation)** An OMG standard for modelling business decisions (Object Management Group 2024). Supported in Camunda 8 as a complement to BPMN.
- End-to-end process** A business process viewed from its triggering event through to its final business outcome — in the BPA Lab, from the arrival of a customer order to the delivery of a bicycle.
- Fischertechnik Industry 4.0 factory** A modular model of an Industry-4.0 factory by Fischertechnik. In the BPA Lab, it serves as the physical demonstration system.
- gRPC** A high-performance Remote Procedure Call framework originated at Google. In the BPA Lab, Camunda 8 uses gRPC to communicate with job workers.
- Happy path** The path through a process in which no errors, exceptions, or edge cases occur.
- IoT (Internet of Things)** The networking of physical devices with the internet and with each other.
- Job worker** In Camunda 8, an external application that picks up service tasks from the workflow and executes them. In the BPA Lab, job workers are implemented in Java.
- MQTT (Message Queuing Telemetry Transport)** A lightweight publish-subscribe protocol widely used in the IoT space. In the BPA Lab, MQTT is the communication protocol between job workers and controllers.
- Openrouteservice** A public API for routing and distance calculation based on OpenStreetMap data. Used in the BPA Lab's shipment process.
- Process mining** Data-driven analysis of business processes based on event logs from information systems. A central teaching and demonstration component in the BPA Lab (Aalst 2016).

B. Glossary

Process application In the BPA Lab, a logical grouping of job workers and process models belonging to one business domain (e.g. *Order Management*).

Production order An order to manufacture a specific product. In the BPA Lab, one production order is created per product line in a customer order.

RPA (Robotic Process Automation) A technology to automate manual, repetitive tasks in existing applications via their user interface.

Workflow engine A software component that executes process models (e.g. expressed in BPMN), manages the state of process instances, and distributes work to humans or job workers.

XES (eXtensible Event Stream) An IEEE standard for event logs in process mining — a format for the exchange of process event data between systems.

C. References

- Aalst, Wil M. P. van der. 2016. *Process Mining: Data Science in Action*. 2nd ed. Springer. <https://doi.org/10.1007/978-3-662-49851-4>.
- Brown, Simon. 2018. *The C4 Model for Visualising Software Architecture*. Leanpub. <https://c4model.com/>.
- Camunda Services GmbH. 2026. *Camunda 8 Documentation*. <https://docs.camunda.io/>.
- Evans, Eric. 2003. *Domain-Driven Design: Tackling Complexity in the Heart of Software*. Addison-Wesley.
- Object Management Group. 2014. *Business Process Model and Notation (BPMN), Version 2.0.2*. <https://www.omg.org/spec/BPMN/>.
- Object Management Group. 2024. *Decision Model and Notation (DMN)*. <https://www.omg.org/spec/DMN/>.

